



Data Handling: Import, Cleaning and Visualisation

Lecture 7:

Data Sources, Data Gathering, Data Import

Prof. Dr. Ulrich Matter

Part II: Data gathering and preparation

Date	Topic
17.11.2022	Data sources, data gathering, data import
24.11.2022	Guest Lecture
24.11.2022	Exercises/Workshop 4: Data gathering, data import
01.12.2022	Data preparation and manipulation

Part III: Analysis, visualisation, output

Date	Topic
08.12.2022	Basic statistics and data analysis with R
08.12.2022	Exercises/Workshop 5: Data preparation and applied data analysis with R
15.12.2022	Visualisation, dynamic documents
21.12.2022	Exercises/Workshop 6: Visualization, dynamic documents
22.12.2022	Summary, Wrap-Up, Q&A, Feedback
22.12.2022	Exam for Exchange Students

Data sources



Or, go to www.menti.com and use the code 58 42 49 4,
or visit the direct link: <https://www.menti.com/al1a9ehfkofo>

Formats/data structures in economics

CSV (typical for rectangular/table-like data)

Variants of CSV (tab-delimited, fix length etc.)

XML and JSON (useful for complex/high-dimensional data sets)

HTML (a markup language to define the structure and layout of webpages)

Unstructured text

Formats/data structures in economics

Excel spreadsheets (`.xls`)

Formats specific to statistical software packages (SPSS: `.sav`, STATA: `.dat`, etc.)

Built-in R datasets

Binary formats

Organize your data pipeline!

One R script to gather/import data.

The beginning of your data pipeline!

A Template/Blueprint

Tell your future self what this script is all about

```
#####  
# Data Handling Course: Example Script for Data Gathering and Import  
#  
# Imports data from ...  
# Input: links to data sources (data comes in ... format)  
# Output: cleaned data as CSV  
#  
# U. Matter, St. Gallen, 2020  
#####
```

Script sections

Recall: programming tasks can often be split into smaller tasks.

Use sections to implement task-by-task and keep order.

In RStudio: Use ----- to indicate the beginning of sections.

Start with a 'meta'-section.

Script sections

```
#####  
# Data Handling Course: Example Script for Data Gathering and Import  
#  
# Imports data from ...  
# Input: links to data sources (data comes in ... format)  
# Output: cleaned data as CSV  
#  
# U. Matter, St. Gallen, 2018  
#####  
  
# SET UP -----  
# load packages  
library(tidyverse)  
  
# set fix variables  
INPUT_PATH <- "/rawdata"  
OUTPUT_FILE <- "/final_data/datafile.csv"
```

Script sections

Finally we add sections with the actual code (in the case of a data import script, maybe one section per data source)

```
#####  
# Project XY: Data Gathering and Import  
#  
# This script is the first part of the data pipeline of project XY.  
# It imports data from ...  
# Input: links to data sources (data comes in ... format)  
# Output: cleaned data as CSV  
#  
# U. Matter, St. Gallen, 2018  
#####  
  
# SET UP -----  
# load packages  
library(tidyverse)  
  
# set fix variables  
INPUT_PATH <- "/rawdata"  
OUTPUT_FILE <- "/final_data/datafile.csv"  
  
# IMPORT RAW DATA FROM CSVs -----
```


Loading built-in datasets

In order to load such datasets, simply use the `data()`-function:

```
data(swiss)
```

Inspect the data after loading

```
# inspect the structure
```

```
str(swiss)
```

```
## 'data.frame':    47 obs. of  6 variables:
## $ Fertility      : num  80.2 83.1 92.5 85.8 76.9 76.1 83.8 92.4 82.4 82.9 ...
## $ Agriculture    : num  17 45.1 39.7 36.5 43.5 35.3 70.2 67.8 53.3 45.2 ...
## $ Examination    : int   15 6 5 12 17 9 16 14 12 16 ...
## $ Education      : int   12 9 5 7 15 7 7 8 7 13 ...
## $ Catholic       : num   9.96 84.84 93.4 33.77 5.16 ...
## $ Infant.Mortality: num  22.2 22.2 20.2 20.3 20.6 26.6 23.6 24.9 21 24.4 ...
```

```
# look at the first few rows
```

```
head(swiss)
```

##	Fertility	Agriculture	Examination	Education	Catholic	Infant.Mortality
## Courtelary	80.2	17.0	15	12	9.96	22.2
## Delemont	83.1	45.1	6	9	84.84	22.2
## Franches-Mnt	92.5	39.7	5	5	93.40	20.2
## Moutier	85.8	36.5	12	7	33.77	20.3
## Neuveville	76.9	43.5	17	15	5.16	20.6
## Porrentruy	76.1	35.3	9	7	90.57	26.6

Comma Separated Values (CSV)

The `swiss`-dataset would look like this when stored in a CSV:

```
"District", "Fertility", "Agriculture", "Examination", "Education", "Catholic", "Infant.Mor  
"Courtelary", 80.2, 17, 15, 12, 9.96, 22.2
```

What do we need to read this format properly?

Parsing CSVs in R

`read.csv()` (basic R distribution)

Returns a `data.frame`

```
swiss_imported <- read.csv("data/swiss.csv")
```

Parsing CSVs in R

Alternative: `read_csv()` (readr/tidyr-package)

Returns a `tibble`.

Used in Wickham and Grolemund (2017).

```
swiss_imported <- read_csv("data/swiss.csv")
```

Import and parsing with `readr`

Why `readr`?

- Functions for all common rectangular data formats.

- Consistent syntax.

- More robust and faster than similar functions in basic R.

Alternative: The `data.table`-package (handling large datasets).

Basic usage of readr functions

Parse the first lines of the swiss dataset directly like this...

```
library(readr)
```

```
read_csv('"District","Fertility","Agriculture","Examination","Education","Catholic","Courtelary",80.2,17,15,12,9.96,22.2')
```

```
## Rows: 1 Columns: 7
```

```
## — Column specification
```

```
## Delimiter: ","
```

```
## chr (1): District
```

```
## dbl (6): Fertility, Agriculture, Examination, Education, Catholic, Infant.Mortality
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
## # A tibble: 1 × 7
```

```
##   District    Fertility Agriculture Examination Education Catholic Infant.Mortality
```

```
##   <chr>         <dbl>         <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
```

```
## 1 Courtelary    80.2           17            15            12            9.96          22.2
```

or read the entire `swiss` dataset by pointing to the file

```
swiss <- read_csv("data/swiss.csv")
```

Basic usage of readr functions

In either case, the result is a `tibble`:

```
swiss
```

```
## # A tibble: 47 × 7
##   District      Fertility Agriculture Examination Education Catholic Infant.Mortality
##   <chr>          <dbl>         <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
## 1 Courtelary      80.2           17           15           12           9.96          2
## 2 Delemont        83.1           45.1          6            9           84.8          2
## 3 Franches-Mnt    92.5           39.7          5            5           93.4          2
## 4 Moutier         85.8           36.5         12            7           33.8          2
## 5 Neuveville      76.9           43.5         17           15           5.16          2
## 6 Porrentruy      76.1           35.3          9            7           90.6          2
## 7 Broye           83.8           70.2         16            7           92.8          2
## 8 Glane           92.4           67.8         14            8           97.2          2
## 9 Gruyere         82.4           53.3         12            7           97.7          2
## 10 Sarine         82.9           45.2         16           13           91.4          2
## # ... with 37 more rows
```

Basic usage of `readr` functions

Other `readr` functions have practically the same syntax and behavior.

`read_tsv()` (tab-separated)

`read_fwf()` (fixed-width)

...

Parsing CSVs

Recognizing columns and rows is one thing...

```
swiss
```

```
## # A tibble: 47 × 7
##   District      Fertility Agriculture Examination Education Catholic Infant.Mortality
##   <chr>         <dbl>         <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
## 1 Courtelary    80.2           17            15            12            9.96          2
## 2 Delemont      83.1           45.1          6             9            84.8          2
## 3 Franches-Mnt  92.5           39.7          5             5            93.4          2
## 4 Moutier       85.8           36.5          12            7            33.8          2
## 5 Neuveville    76.9           43.5          17            15            5.16          2
## 6 Porrentruy    76.1           35.3          9             7            90.6          2
## 7 Broye         83.8           70.2          16            7            92.8          2
## 8 Glane         92.4           67.8          14            8            97.2          2
## 9 Gruyere       82.4           53.3          12            7            97.7          2
## 10 Sarine       82.9           45.2          16            13           91.4          2
## # ... with 37 more rows
```

What else did `read_csv()` recognize?

Parsing CSVs

Recall the introduction to data structures and data types in R

How does R represent data in RAM

Structure: `data.frame/tibble`, etc.

Types: `character, numeric`, etc.

Parsers in `read_csv()` guess the data **types**.

Parsing CSV-columns

`"12:00": type character?`

Parsing CSV-columns

`"12:00": type character?`

What about `c("12:00", "midnight", "noon")`?

Parsing CSV-columns

`"12:00": type character?`

What about `c("12:00", "midnight", "noon")`?

And now `c("12:00", "14:30", "20:01")`?

Parsing CSV-columns

Let's test it!

```
read_csv('A,B
         12:00, 12:00
         14:30, midnight
         20:01, noon')
```

```
## Rows: 3 Columns: 2
## — Column specification —————
## Delimiter: ",",
## chr (1): B
## time (1): A
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
## # A tibble: 3 × 2
##   A      B
##   <time> <chr>
## 1 12:00  12:00
## 2 14:30  midnight
## 3 20:01  noon
```

Parsing CSV-columns

Let's test it!

```
read_csv('A,B
         12:00, 12:00
         14:30, midnight
         20:01, noon')
```

```
## Rows: 3 Columns: 2
## — Column specification —————
## Delimiter: ","
## chr  (1): B
## time (1): A
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
## # A tibble: 3 × 2
##   A      B
##   <time> <chr>
## 1 12:00  12:00
## 2 14:30  midnight
## 3 20:01  noon
```

How can `read_csv()` distinguish the two cases?

Parsing CSV-columns: guess types

Under the hood `read_csv()` used the `guess_parser()` - function to determine which type the two vectors likely contain:

```
guess_parser(c("12:00", "midnight", "noon"))
```

```
## [1] "character"
```

```
guess_parser(c("12:00", "14:30", "20:01"))
```

```
## [1] "time"
```


Spreadsheets/Excel

Needs additional R-package: `readxl`.

```
# install the package  
install.packages("readxl")
```

Spreadsheets/Excel

Then we load this additional package ('library') and use the package's `read_excel()`-function to import data from an excel-sheet.

```
# load the package  
library(readxl)  
  
# import data from a spreadsheet  
swiss_imported <- read_excel("data/swiss.xlsx")
```

Data from other data analysis software

STATA, SPSS, etc.

Additional packages needed:

`foreign`

`haven`

Parsers (functions) for many foreign formats.

For example, `read_spss()` for SPSS' `.sav`-format.

Data from other data analysis software

```
# install the package (if not yet installed):  
# install.packages("haven")
```

```
# load the package  
library(haven)
```

```
# read the data  
swiss_imported <- read_spss("data/swiss.sav")
```


XML in R

```
## {xml_document}  
## <customers>  
## [1] <person>\n  <name>John Doe</name>\n  <orders>\n    <product> x </product>\n  </orders>\n## [2] <person>\n  <name>Peter Pan</name>\n  <orders>\n    <product> a </product>\n  </orders>\n</customers>
```

load packages

```
library(xml2)
```

parse XML, represent XML document as R object

```
xml_doc <- read_xml("data/customers.xml")  
xml_doc
```

XML in R: tree-structure

'customers' is the root-node, 'persons' are it's children:

```
# navigate downwards
```

```
persons <- xml_children(xml_doc)
persons
```

```
## {xml_nodeset (2)}
```

```
## [1] <person>\n  <name>John Doe</name>\n  <orders>\n    <product> x </product>\n  </orders>\n  </person>\n## [2] <person>\n  <name>Peter Pan</name>\n  <orders>\n    <product> a </product>\n  </orders>\n  </person>
```

XML in R: tree-structure

Navigate sideways and upwards

```
# navigate sideways
```

```
persons[1]
```

```
## {xml_nodeset (1)}
```

```
## [1] <person>\n  <name>John Doe</name>\n  <orders>\n    <product> x </product>\n
```

```
xml_siblings(persons[[1]])
```

```
## {xml_nodeset (1)}
```

```
## [1] <person>\n  <name>Peter Pan</name>\n  <orders>\n    <product> a </product>\n
```

```
# navigate upwards
```

```
xml_parents(persons)
```

```
## {xml_nodeset (1)}
```

```
## [1] <customers>\n  <person>\n    <name>John Doe</name>\n    <orders>\n      <product>
```

XML in R: tree-structure

Extract specific parts of the data:

```
# find data via XPath  
customer_names <- xml_find_all(xml_doc, xpath = ".//name")  
# extract the data as text  
xml_text(customer_names)
```

```
## [1] "John Doe" "Peter Pan"
```

JSON in R

```
# load packages
library(jsonlite)

# parse the JSON-document shown in the example above
json_doc <- fromJSON("data/person.json")

# look at the structure of the document
str(json_doc)

## List of 6
## $ firstName : chr "John"
## $ lastName  : chr "Smith"
## $ age       : int 25
## $ address   :List of 4
## ..$ streetAddress: chr "21 2nd Street"
## ..$ city          : chr "New York"
## ..$ state         : chr "NY"
## ..$ postalCode    : chr "10021"
## $ phoneNumber:'data.frame': 2 obs. of 2 variables:
## ..$ type : chr [1:2] "home" "fax"
## ..$ number: chr [1:2] "212 555-1234" "646 555-4567"
## $ gender    :List of 1
## ..$ type: chr "male"
```

JSON in R

The nesting structure is represented as a **nested list**:

```
# navigate the nested lists, extract data  
# extract the address part  
json_doc$address
```

```
## $streetAddress  
## [1] "21 2nd Street"  
##  
## $city  
## [1] "New York"  
##  
## $state  
## [1] "NY"  
##  
## $postalCode  
## [1] "10021"
```

```
# extract the gender (type)  
json_doc$gender$type
```

```
## [1] "male"
```


Recognize the problem

```
FILE <- "../..data/hastamanana.txt"  
hasta <- readLines(FILE)  
hasta
```

```
## [1] "Hasta Ma\xflana!"
```

(`readLines()` simply reads the content of a text file line by line.)

Guess encoding

Recall that there are no meta data in csv or plain text file informing you about the encoding.

If no other information is available, we need to make an educated guess.

`readr` provides a function that does just that: `guess_encoding()`

```
guess_encoding(FILE)
```

```
## # A tibble: 3 × 2
##   encoding    confidence
##   <chr>         <dbl>
## 1 ISO-8859-2     0.64
## 2 ISO-8859-1     0.42
## 3 ISO-8859-9     0.21
```

Handling encoding issues

`iconv()`: convert a character vector from one encoding to another encoding.

Use the guessed encoding for the `from` argument

```
iconv(hasta, from = "ISO-8859-2", to = "UTF-8")
```

```
## [1] "Hasta Mañana!"
```

```
iconv(hasta, from = "ISO-8859-1", to = "UTF-8")
```

```
## [1] "Hasta Mañana!"
```


References

Wickham, Hadley, and Garrett Grolemund. 2017. Sebastopol, CA: O'Reilly. <http://r4ds.had.co.nz/>.